

# Document DB (On Azure)

Dan Matthews

# Overview

- ▶ Schema-free
- ▶ Storage
- ▶ Indexing Policies
- ▶ Sharding

# Azure Storage Options

- ▶ Azure SQL
- ▶ DocumentDB
- ▶ Azure Tables

# Meet Azure DocumentDB

- ▶ Azure DocumentDB is a NoSQL document database service designed from the ground up to natively support JSON and JavaScript directly inside the database engine.
- ▶ It is great for applications that run in the cloud, when predictable throughput, low latency, and flexible queries.

# Schema-free JSON data

- ▶ A common problem for developers is that application schemas constantly evolve.
- ▶ Schema: A database schema is a structure in a formal language supported by the database management system that refers to the organization of data as a blue print of how a database is constructed.
  - ▶ Divided into database tables in the case of relational databases
- ▶ DocumentDB automatically indexes all JSON documents added to the database, then lets you use familiar SQL syntax to query them without specifying schema or secondary indices up front.

# Benefits of schema-free

- ▶ One benefit is that you can often represent in a single entity a construct that would require several tables to properly represent in a relational db.
- ▶ No schema migrations: Since DocumentDB is schema free, your code will define your schema
- ▶ Also, there is a clear path to horizontal scalability which can improve the speed of queries on large data sets
- ▶ Even Hadoop jobs over data stored in DocumentDB are supported using the connector

# Storage

- ▶ DocumentDB stores its data as a native JSON Object
  - ▶ This allows attachments to documents, while azure tables (Azures other NoSQL storage option) stores data in XML.
  - ▶ The advantage of storing native JSON objects is that the data maps directly into JSON objects in the client application code
- ▶ DocumentDB can store petabytes of information.
  - ▶ One petabyte is approximately  $2^{50}$  bytes, or 2,000 terabytes

# DocumentDB Indexing Policies

- ▶ DocumentDB is a true schema-free database. It does not assume or require any schema for the JSON documents it indexes.
  - ▶ This allows you to quickly define and iterate on application data models
- ▶ As you add documents to a collection, DocumentDB automatically indexes all document properties so they are available for you to query.
  - ▶ Automatic indexing also allows you to store heterogeneous types of documents.
  - ▶ DocumentDB supports a sustained volume of fast writes while still serving consistent queries.



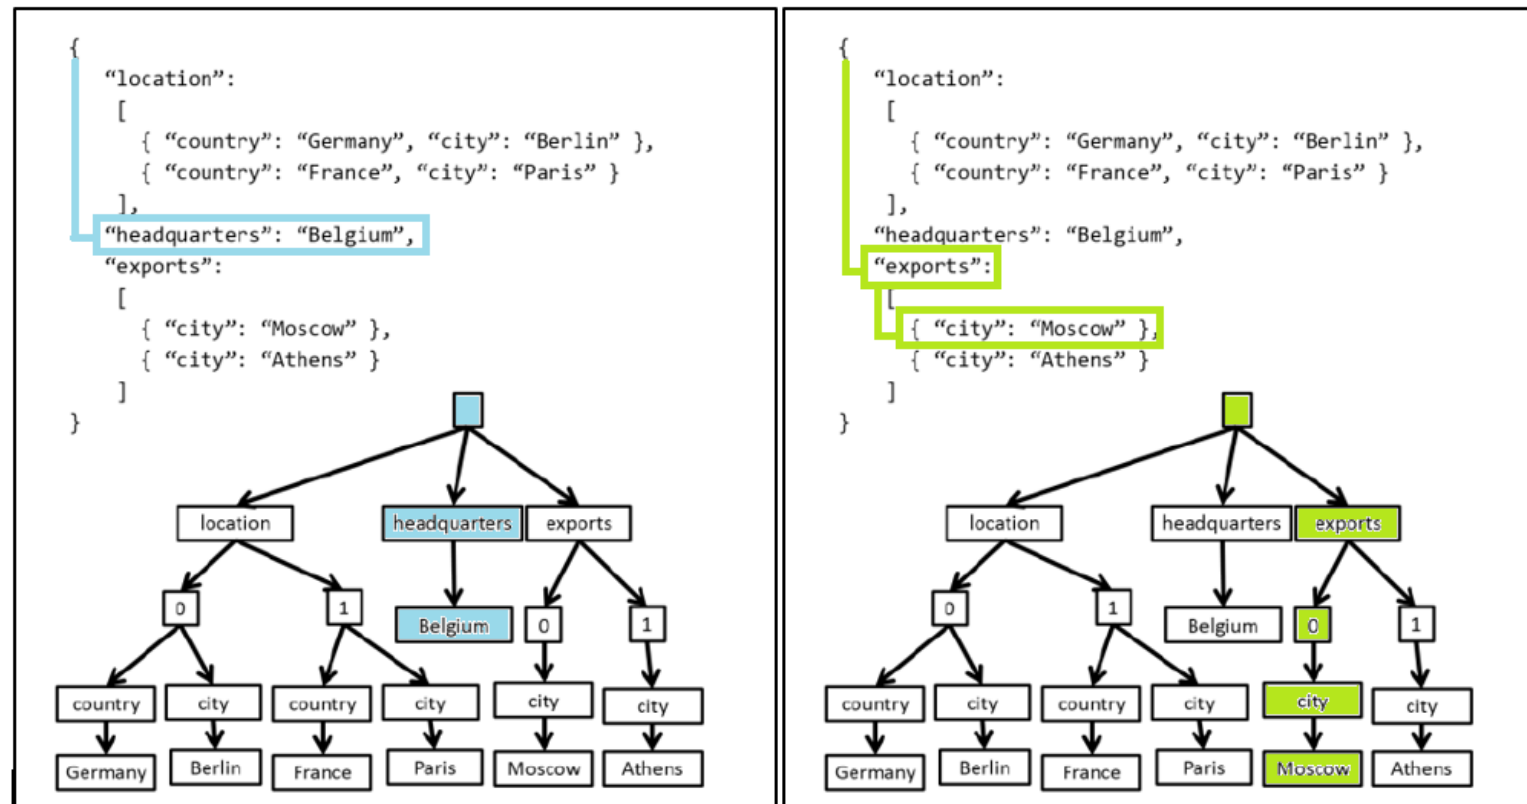
# DocumentDB Indexing Policies(1)

- ▶ The DocumentDB indexing subsystem is designed to support:
  - ▶ Efficient, rich hierarchical and relational queries without any schema or index definitions.
  - ▶ Consistent query results while handling a sustained volume of writes. For high write throughput workloads with consistent queries, the index is updated incrementally, efficiently, and online while handling a sustained volume of writes.
  - ▶ Storage efficiency. For cost effectiveness, the on-disk storage overhead of the index is bounded and predictable.
  - ▶ Storage efficiency. For cost effectiveness, the on-disk storage overhead of the index is bounded and predictable.

# DocumentDB Indexing Policies(2)

- ▶ How DocumentDB Indexing works
  - ▶ The indexing in DocumentDB takes advantage of the fact that JSON grammar allows documents to be represented as trees.
  - ▶ For a JSON document to be represented as a tree, a dummy root node needs to be created which parents the rest of the actual nodes in the document underneath
  - ▶ Each label including the array indices in a JSON document becomes a node of the tree.

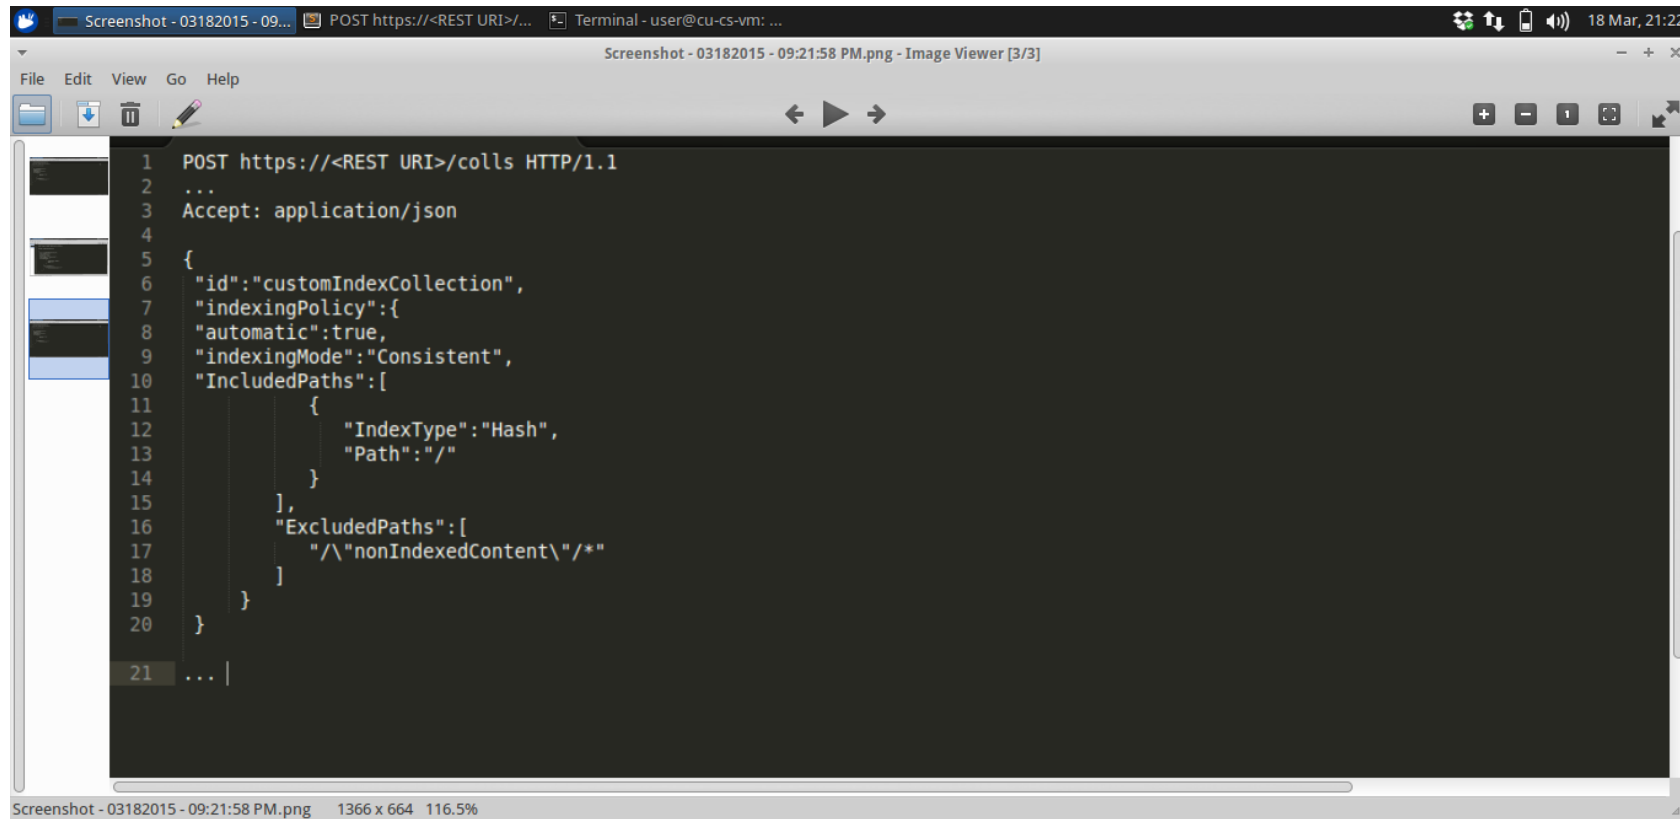
# DocumentDB Indexing Policies(3)



# DocumentDB Indexing Policies(4)

- ▶ In the previous picture: the JSON property `{"headquarters": "Belgium"}` property in the above example corresponds to the path `/"headquarters"/"Belgium"`
- ▶ The JSON array `{"exports": [{"city": "Moscow"}, {"city": "Athens"}]}` correspond to the paths `/"exports"/0/"city"/"Moscow"` and `/"exports"/1/"city"/"Athens"`
  - ▶ **Note** The path representation blurs the boundary between the structure/schema and the instance values in documents, allowing DocumentDB to be truly schema-free.

# Configuring the Indexing Policy of a Collection



The screenshot shows a REST client window with a menu bar (File, Edit, View, Go, Help) and a toolbar. The main area displays a POST request to `https://<REST URI>/colls` with HTTP/1.1. The request body is a JSON object that configures a custom index collection. The JSON includes an `id` of `customIndexCollection`, an `indexingPolicy` object with `automatic` set to `true` and `indexingMode` set to `Consistent`. It also lists `IncludedPaths` (a single path `/` with `IndexType` `Hash`) and `ExcludedPaths` (a single path `/\nonIndexedContent\/*`). The status bar at the bottom indicates the file is `Screenshot - 03182015 - 09:21:58 PM.png` with dimensions `1366 x 664` and zoom `116.5%`.

```
1 POST https://<REST URI>/colls HTTP/1.1
2 ...
3 Accept: application/json
4
5 {
6   "id": "customIndexCollection",
7   "indexingPolicy": {
8     "automatic": true,
9     "indexingMode": "Consistent",
10    "IncludedPaths": [
11      {
12        "IndexType": "Hash",
13        "Path": "/"
14      }
15    ],
16    "ExcludedPaths": [
17      {
18        "Path": "/\nonIndexedContent\/*"
19      }
20    ]
21  }
22 }
```

# Configuring the Indexing Policy of a Collection(2)

- ▶ The Previous sample shows how to set a custom indexing policy during the creation of a collection, by using the DocumentDB REST API. The sample shows the indexing policy expressed in terms of paths, index types, and precisions.
  - ▶ The indexing policy is set to: Automatic
  - ▶ The Indexing mode is set to: Consistent
  - ▶ The indexing type is set to: Hash
- ▶ **Note:** The indexing policy of a collection must be specified at the time of creation. Modifying the indexing policy after collection creation is not allowed, but will be supported in a future release of DocumentDB.
- ▶ **Note:** By default, DocumentDB indexes all paths within documents consistently with a hash index. The internal Timestamp (`_ts`) path is stored with a range index.

# Automatic Indexing

- ▶ You can choose if you want the collection to automatically index all documents or not
- ▶ By default, all documents are automatically indexed, but you can choose to turn it off
- ▶ When indexing is turned off, documents can be accessed only through their self-links or by queries using ID
  - ▶ With automatic indexing turned off, you can still selectively add only specific documents to the index
  - ▶ Conversely, you can leave automatic indexing on and selectively choose to exclude only specific document
- ▶ You can configure the default policy by specifying the value for the automatic property to be true or false.

# Indexing Modes


- ▶ You can choose between synchronous (**Consistent**) and asynchronous (**Lazy**) index updates
- ▶ By default, the index is updated synchronously on each insertion, replacement, or deletion action
  - ▶ This enables the queries to honor the same consistency level as that of the document reads without any delay for the index to catch up.
- ▶ You can also configure certain collections to update their index lazily.
  - ▶ Lazy indexing is great for scenarios where data is written in bursts, and you want to amortize the work required to index content over a longer period of time



# Index Types and Precision

- ▶ There are two supported kinds of index types: Hash and Range
- ▶ Hash:
  - ▶ Choosing an index type of **Hash** enables efficient equality queries. For most use cases, hash indexes do not need a higher precision than the default value of 3 bytes.
- ▶ Range:
  - ▶ Choosing an index type of **Range** enables range queries (using `>`, `<`, `>=`, `<=`, `!=`).
  - ▶ For paths that have large ranges of values, it is recommended to use a higher precision like 6 bytes

# Sharding with DocumentDB



# Sharding with DocumentDB

- ▶ You can achieve a near infinite scale (in terms of storage and throughput) through DocumentDB by horizontally partitioning data, or sharding
- ▶ 3 common sharding patterns:
  - ▶ Range Partitioning
  - ▶ Lookup Partitioning
  - ▶ Hash Partitioning

# Range Partitioning

- ▶ Partitions are assigned based on whether the partitioning key is inside a certain range.
- ▶ An example could be to partition data by timestamp or geography (e.g. zip code is between 30000 and 39999)

# Lookup Partitioning

- ▶ Partitions are assigned based on a lookup directory of discrete values that map to a partition
- ▶ This is generally implemented by creating a lookup map that keeps track of which data is stored on which partition
  - ▶ An example could be to partition data by user

# Hash Partitioning

- ▶ Partitions are based on the value of a hash function - allowing you to evenly distribute across n number of partitions
- ▶ An example could be to partition data by the hash code % 3 of the tenant to evenly distribute tenants across 3 partitions