

...An Introduction

WHY?

- Volume - so much
- Velocity - so fast
- Variety - so many types







The old way just wasn't cutting it


SO HOW DO WE DO IT?

- HDFS
- Map Reduce






It just doesn't scale

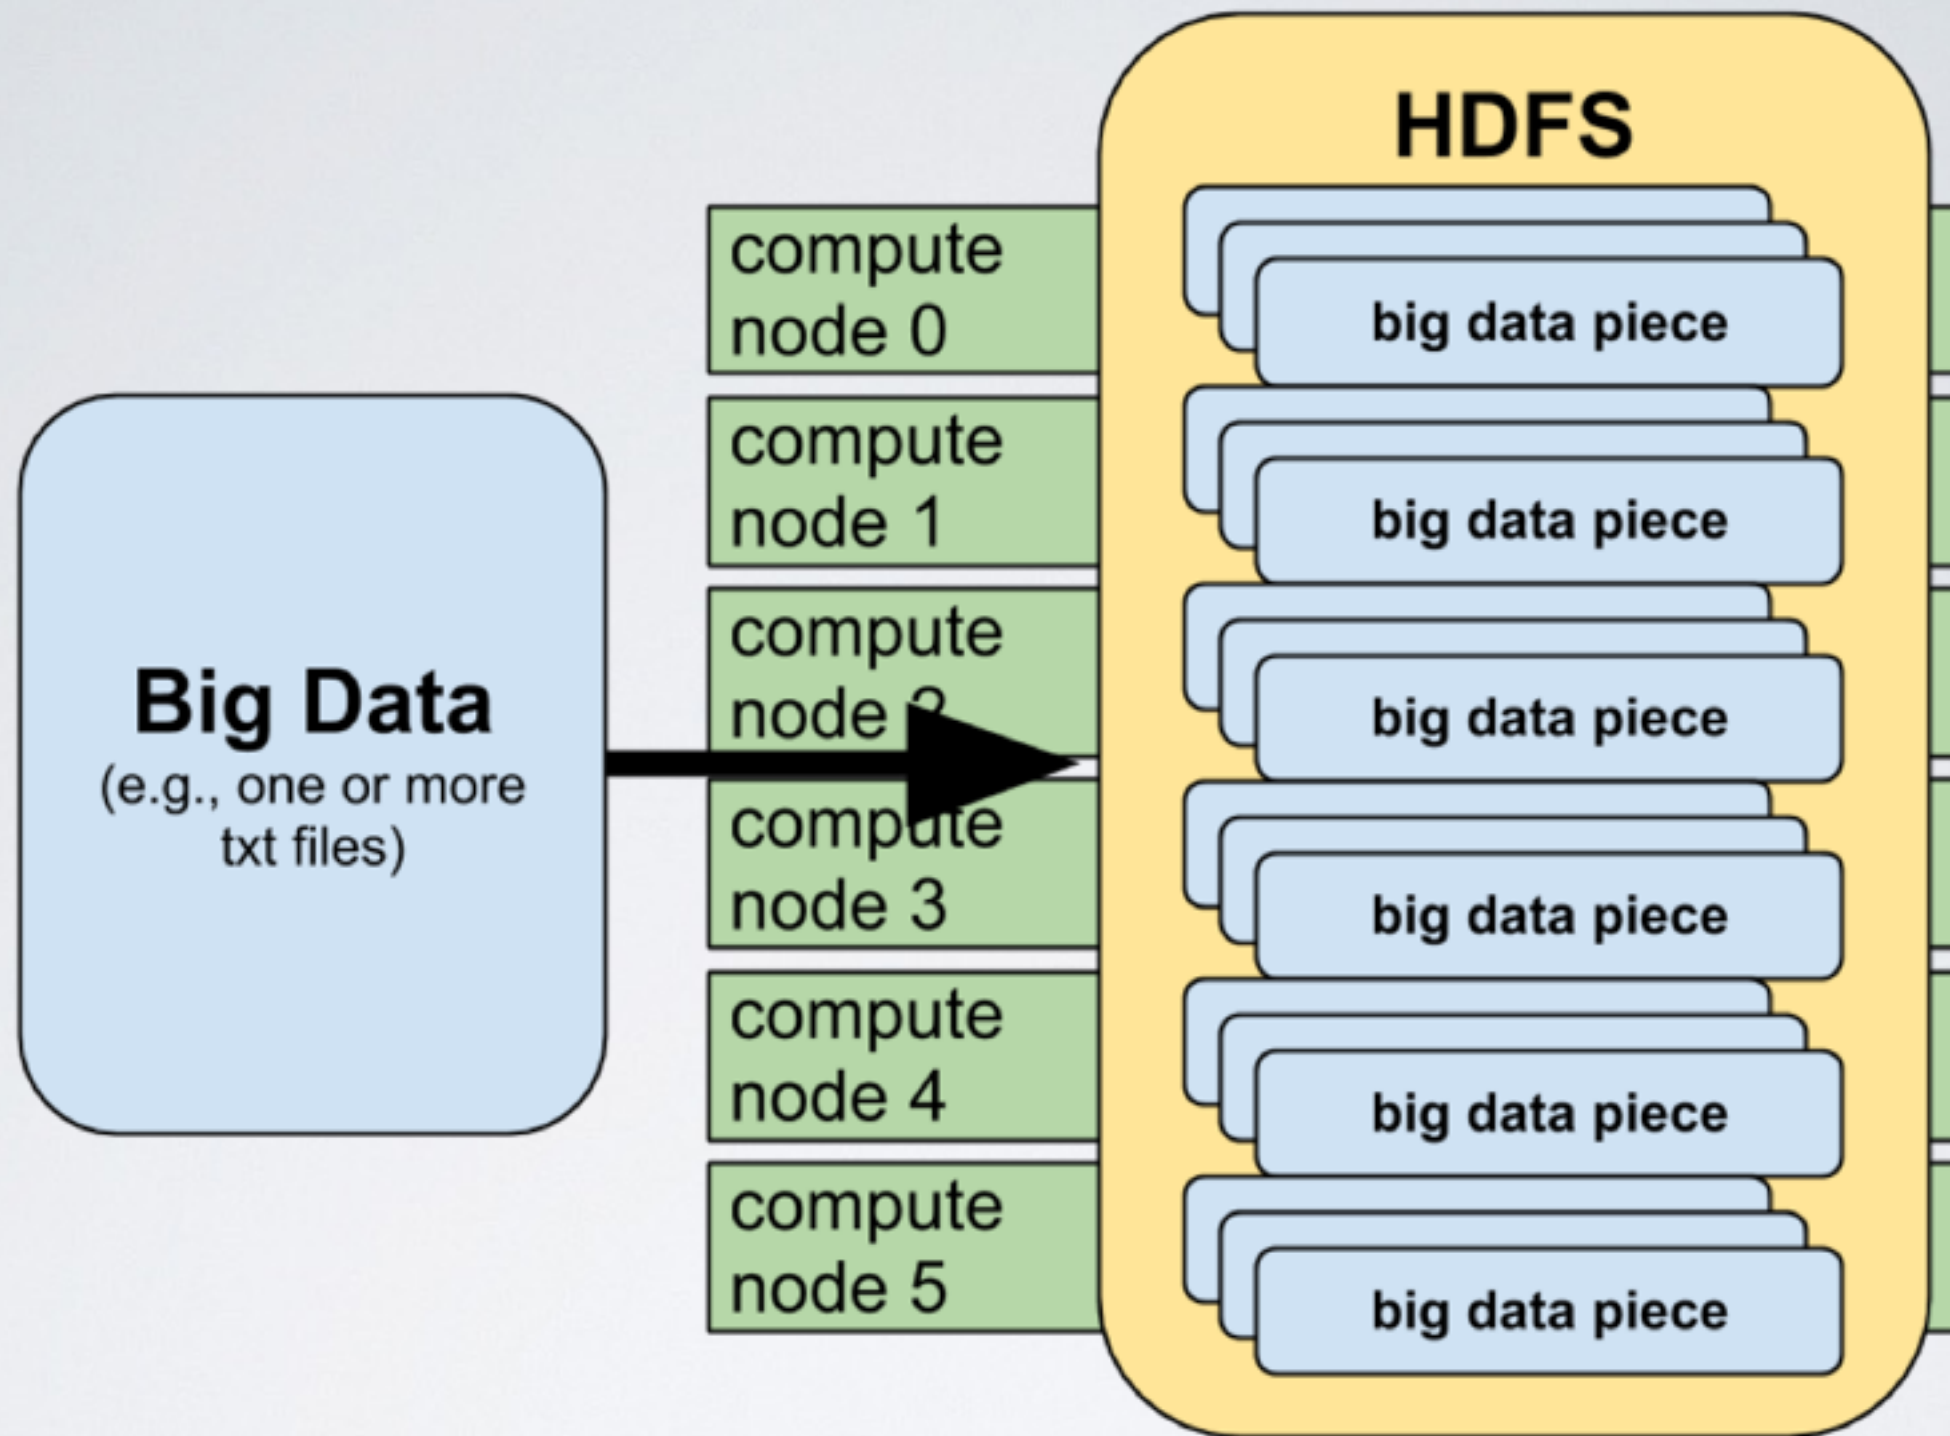


That's better

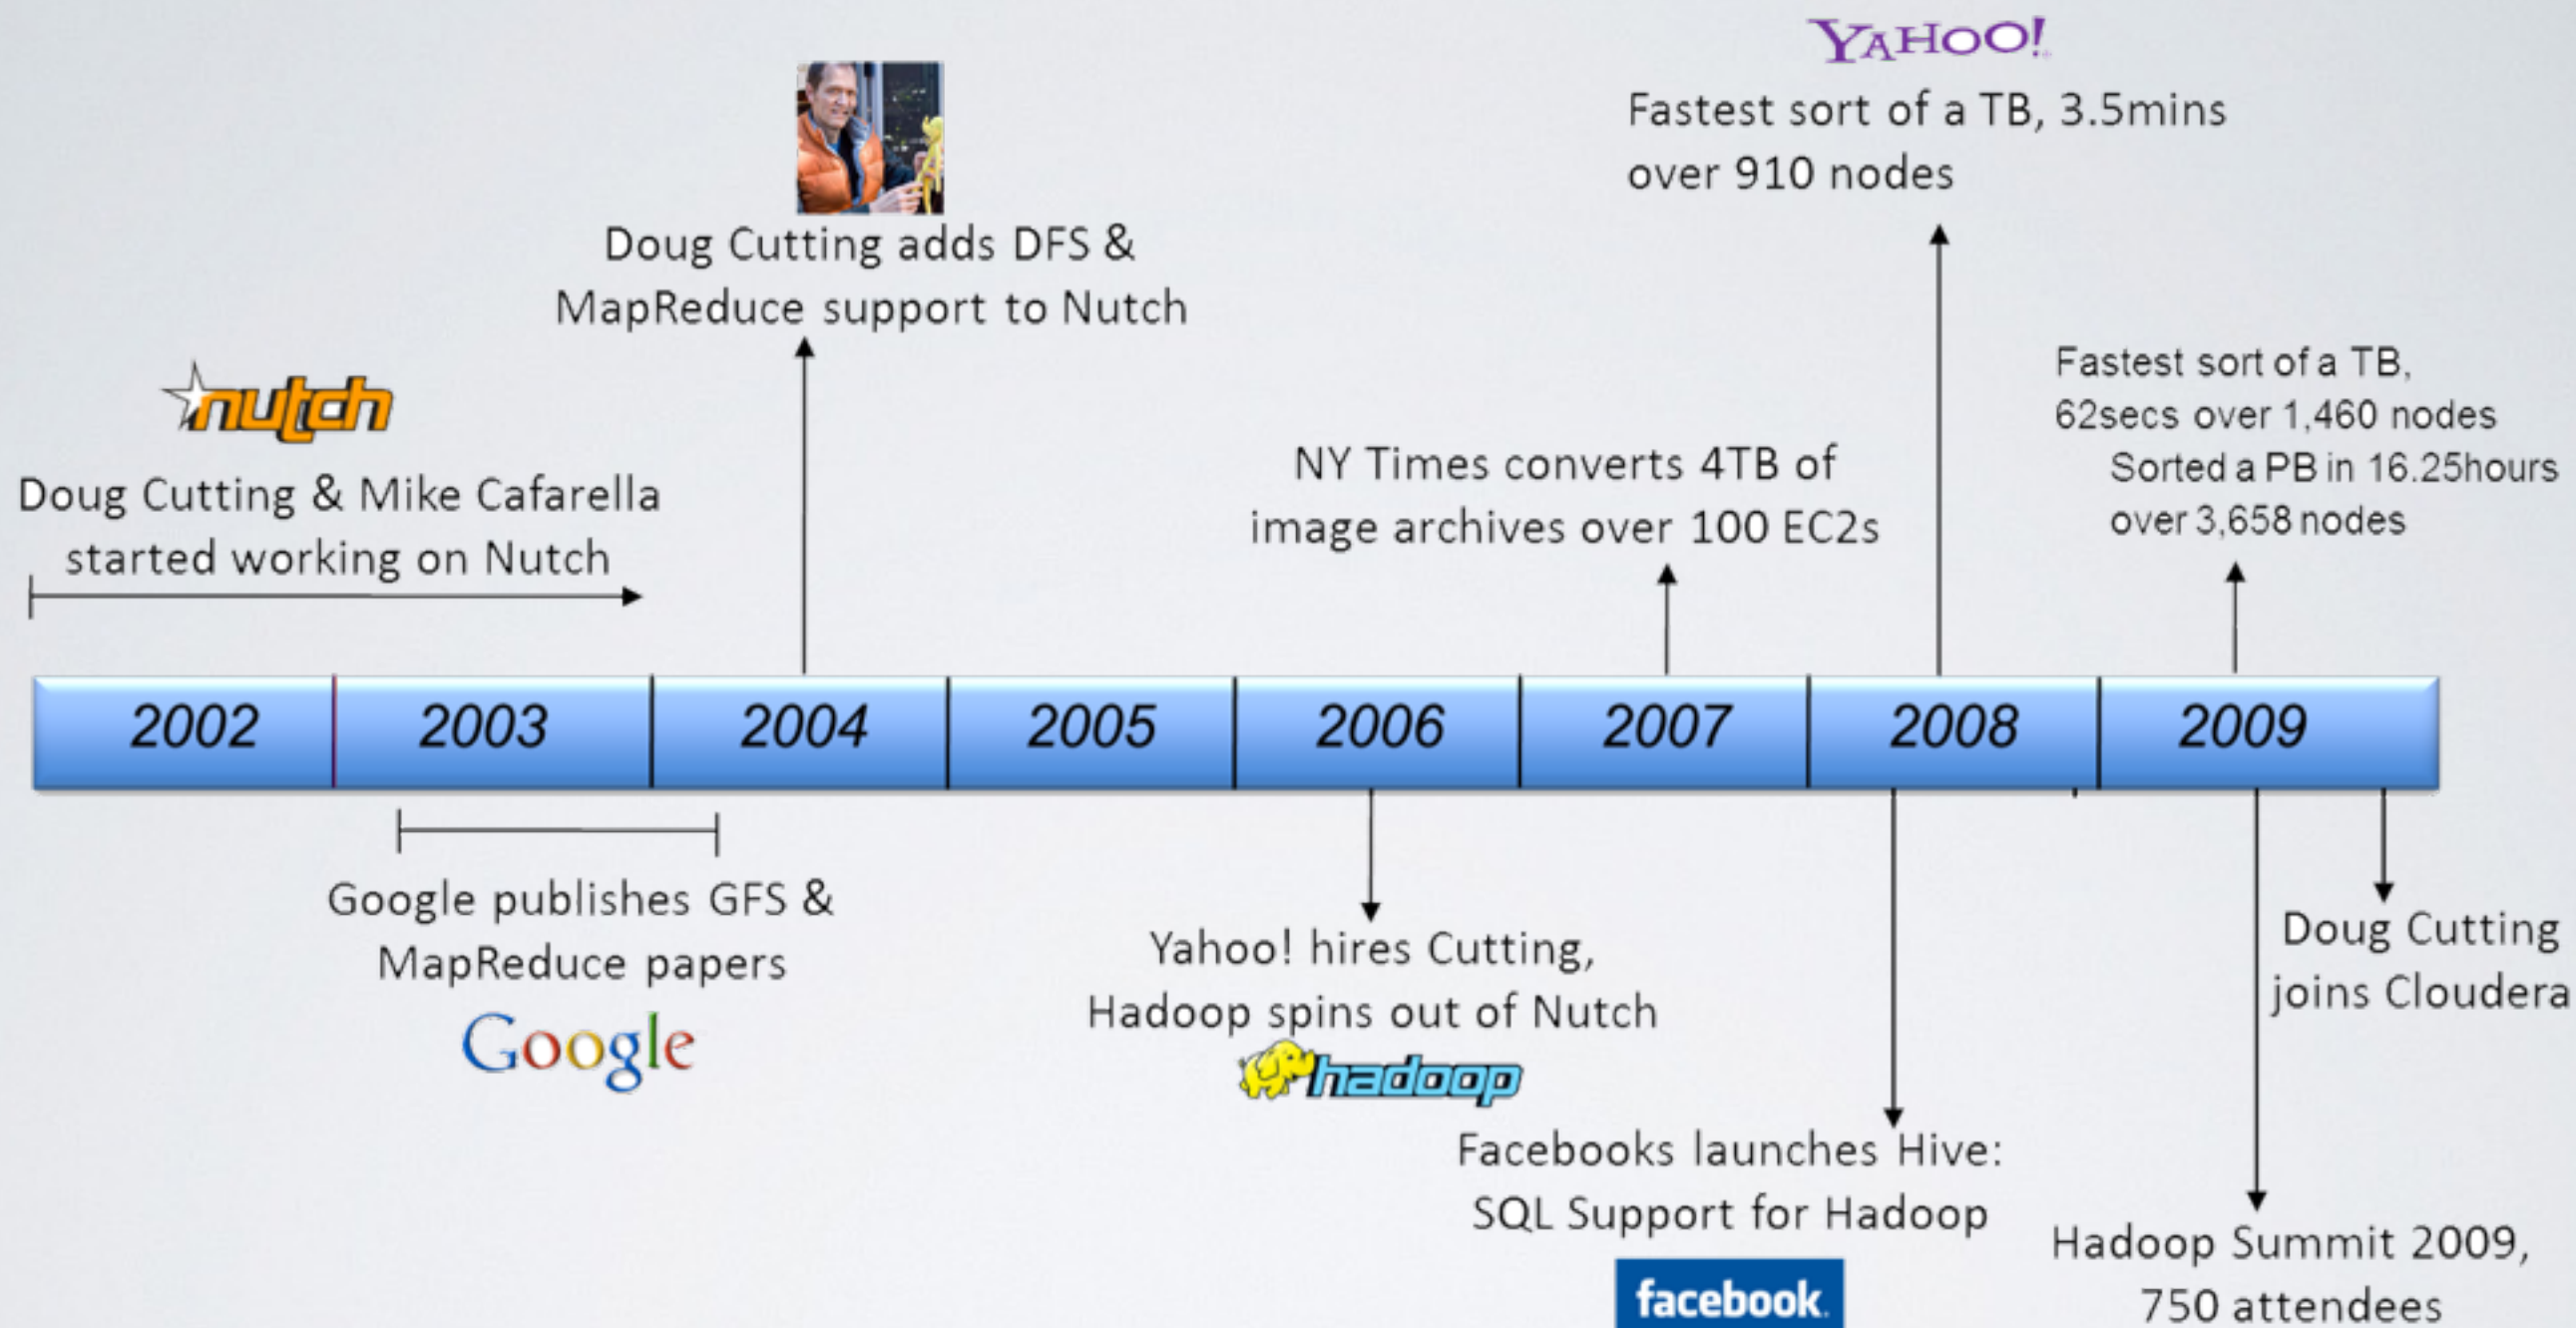
ENTER HADOOP

- Open source Java
- Framework for big data
- Scalable
- Fault tolerant
- Commodity hardware







Distribute the data and the processing



A brief history



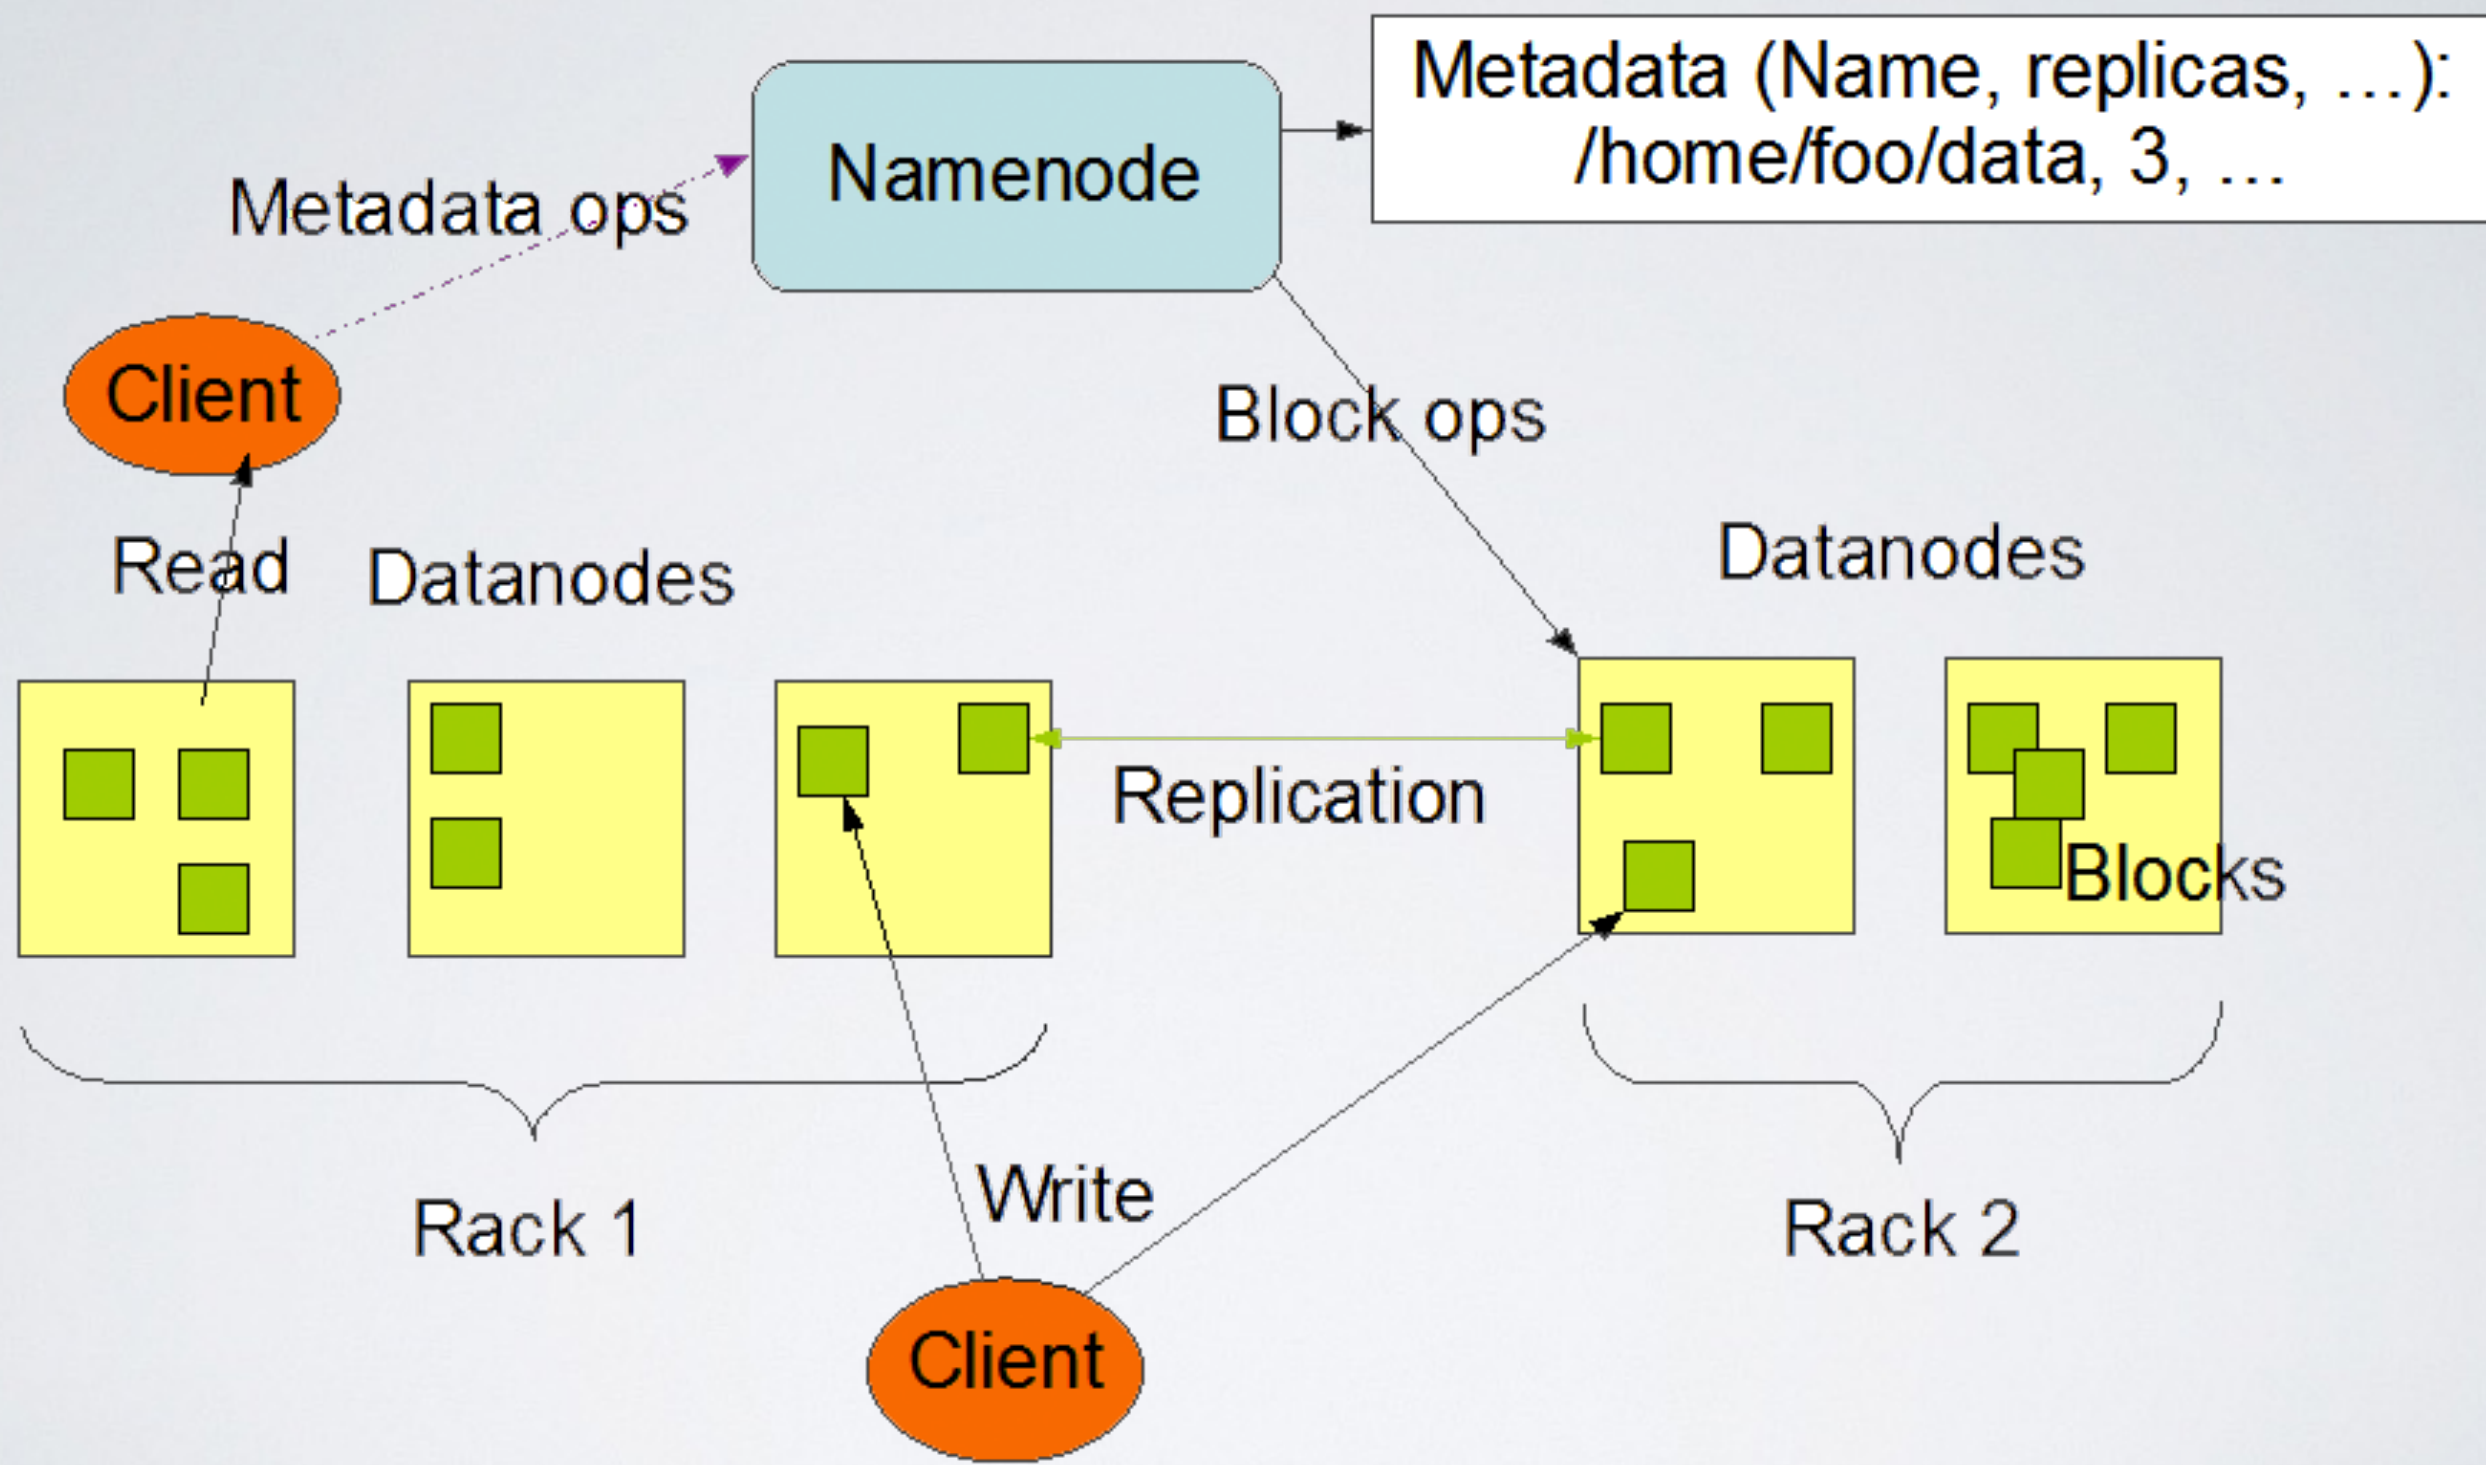
Some hadoop users

HDFS

- Master/slave architecture
- Fault tolerant
- Commodity hardware
- Streaming access to FS data




HDFS Architecture



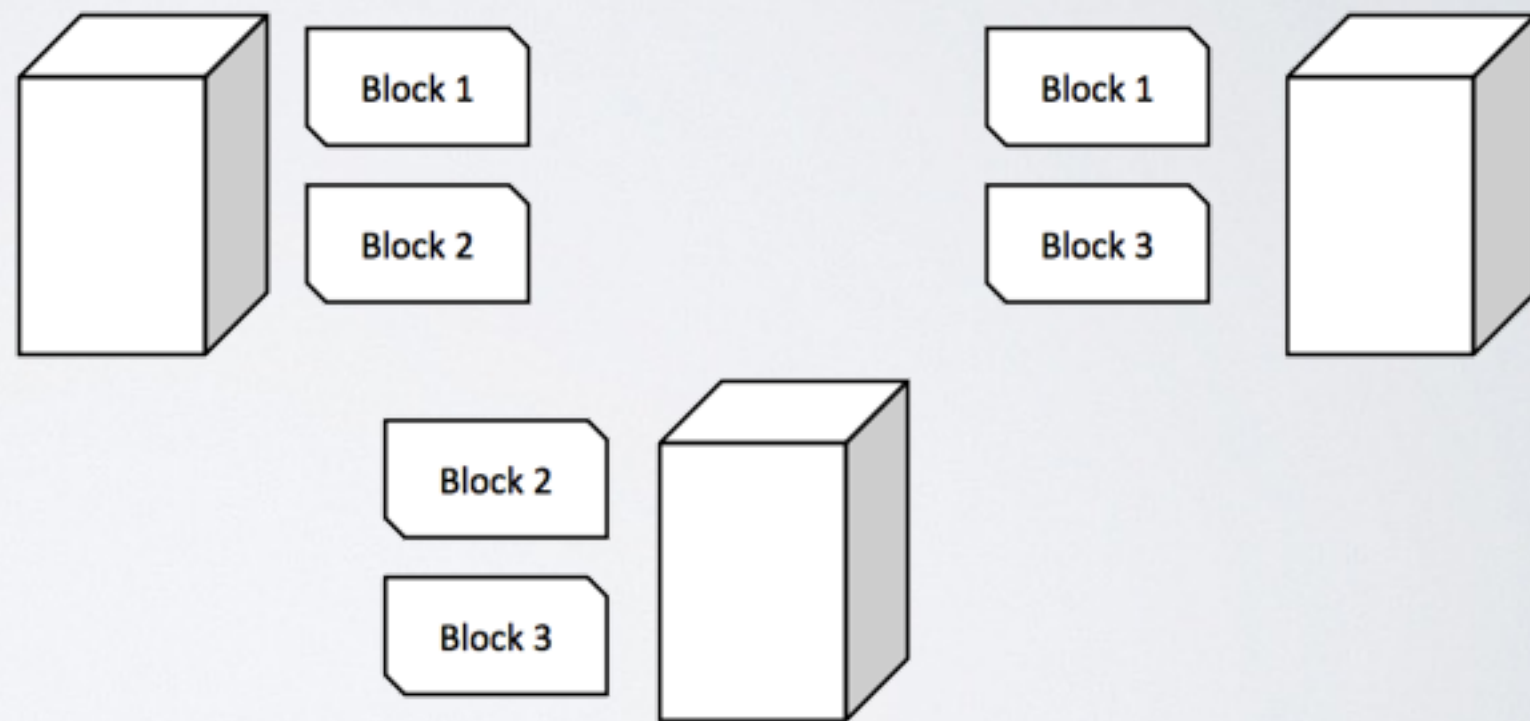
HDFS - BLOCKS

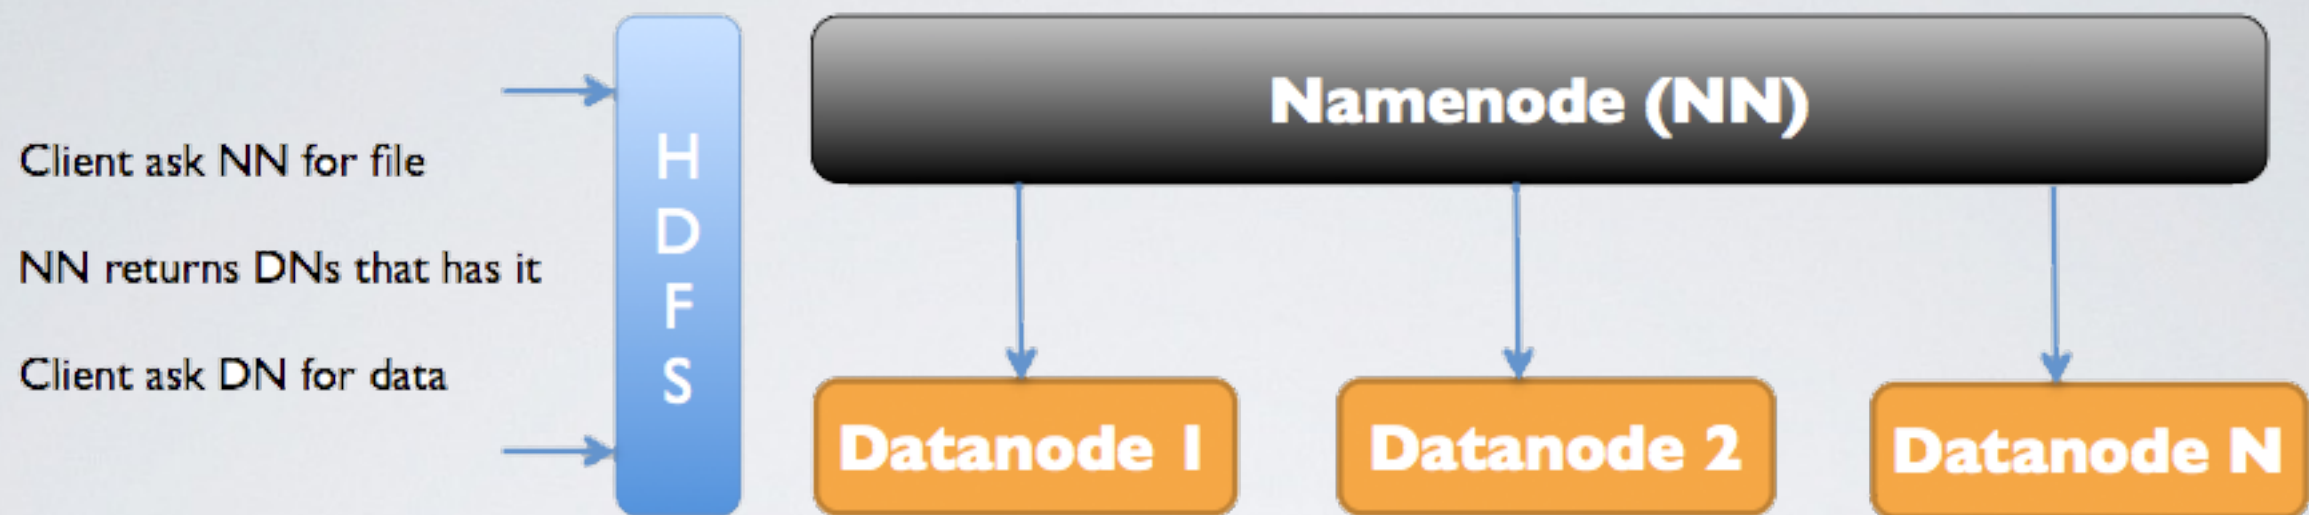
- Splits data into blocks
- Easier to distribute work



HDFS - REPLICATION

- Blocks replicated across nodes
- Fault tolerant
- 3x replications by default



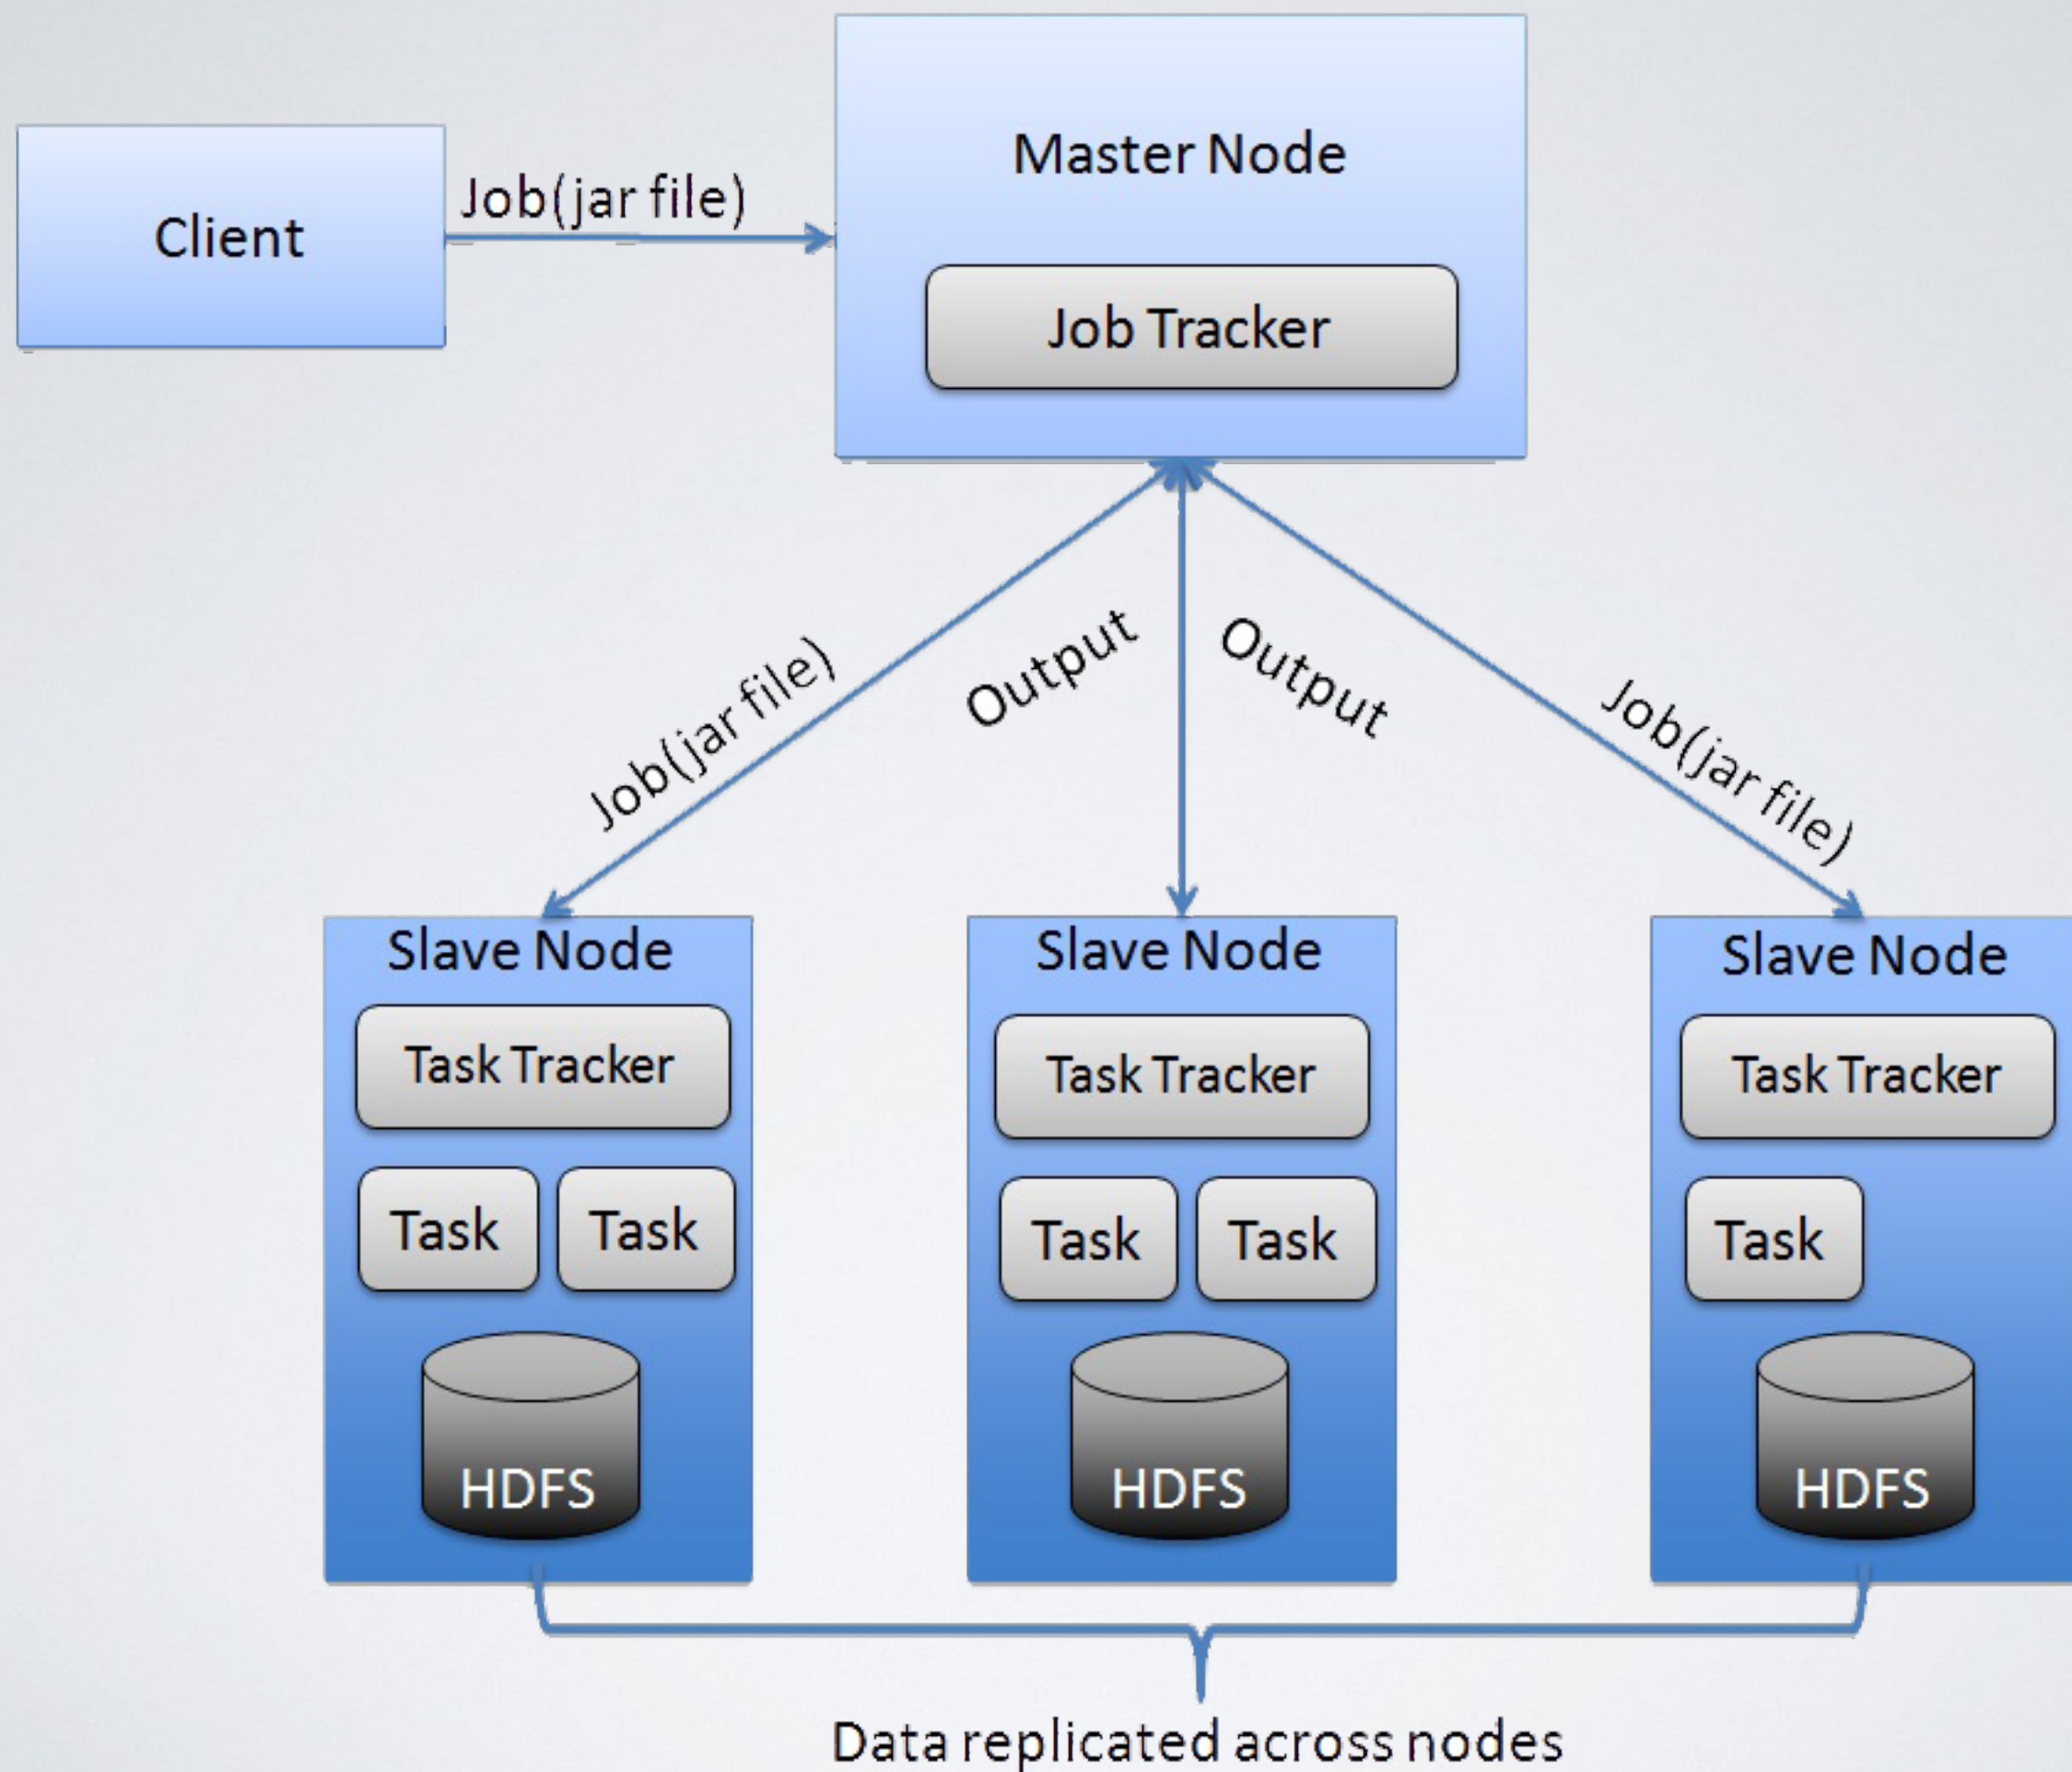


Namenode - Master

- FS metadata
- Block replication
- R/W access to files

Datanode - Slave

- Contains data
- Informs namenode



Map

```
def sqr(n):  
    return n * n
```

```
list = [1,2,3,4]
```

```
map(sqr, list) -> [1,4,9,16]
```

Reduce

```
def add(i, j):  
    return i + j
```

```
list = [1,2,3,4]
```

```
reduce(add, list) -> 10
```

MapReduce

```
def MapReduce(data, mapper, reducer):  
    return reduce(reducer, map(mapper, data))
```

```
MapReduce(list, sqr, add) -> 30
```

```
public static class Map extends MapReduceBase implements Mapper<LongWritable, Text, Text,
IntWritable> {
    private final static IntWritable one = new IntWritable(1);
    private Text word = new Text();

    public void map(LongWritable key, Text value, OutputCollector<Text, IntWritable> output,
Reporter reporter) throws IOException {
        String line = value.toString();
        StringTokenizer tokenizer = new StringTokenizer(line);
        while (tokenizer.hasMoreTokens()) {
            word.set(tokenizer.nextToken());
            output.collect(word, one);
        }
    }
}
```

MAP

wordcount.java

```
public static class Reduce extends MapReduceBase implements Reducer<Text, IntWritable, Text, IntWritable> {  
    public void reduce(Text key, Iterator<IntWritable> values, OutputCollector<Text, IntWritable> output, Reporter reporter) throws IOException {  
        int sum = 0;  
        while (values.hasNext()) {  
            sum += values.next().get();  
        }  
        output.collect(key, new IntWritable(sum));  
    }  
}
```

REDUCE

wordcount.java


```
$ mkdir wordcount_classes
$ javac -classpath ${HADOOP_HOME}/hadoop-${HADOOP_VERSION}-core.jar -d wordcount_classes
WordCount.java
$ jar -cvf /usr/joe/wordcount.jar -C wordcount_classes/ .
```

Assuming that:

- /usr/joe/wordcount/input - input directory in HDFS
- /usr/joe/wordcount/output - output directory in HDFS

Sample text-files as input:

```
$ bin/hadoop dfs -ls /usr/joe/wordcount/input/
/usr/joe/wordcount/input/file01
/usr/joe/wordcount/input/file02
```

```
$ bin/hadoop dfs -cat /usr/joe/wordcount/input/file01
Hello World Bye World
```

```
$ bin/hadoop dfs -cat /usr/joe/wordcount/input/file02
Hello Hadoop Goodbye Hadoop
```

Run the application:

```
$ bin/hadoop jar /usr/joe/wordcount.jar org.myorg.WordCount /usr/joe/wordcount/input
/usr/joe/wordcount/output
```

Output:

```
$ bin/hadoop dfs -cat /usr/joe/wordcount/output/part-00000
Bye 1
Goodbye 1
Hadoop 2
Hello 2
World 2
```

RUNNING IT

wordcount.java