

# Basic Machine Learning

# What is Machine Learning?

- Arthur Samuel (1959). Machine Learning: Field of study that gives computers the ability to learn without being explicitly programmed.
- Tom Mitchell (1998) Well-posed Learning Problem: A computer program is said to learn from experience  $E$  with respect to some task  $T$  and some performance measure  $P$ , if its performance on  $T$ , as measured by  $P$ , improves with experience  $E$ .

# Example Applications

- Spam Filters
- Handwriting and speech recognition
- Product Recommendations (Netflix, Amazon, etc.)
- Computer Vision
- Data mining of large databases in fields like biology, medicine, engineering, etc.

# Types of Machine Learning

- Thee Major Types
  - Supervised Learning: Algorithm is trained using data where the desired output for a given input is known (labeled data).
  - Unsupervised Learning: Algorithm is trained on unlabeled data and left to determine structure on its own.
  - Reinforced Learning: Algorithm is trained by interaction with dynamic environment.

# Supervised Learning

- Arguable the simplest (relatively) paradigm and the focus of this presentation.
- Two major uses:
  - Regression: Measure relationships between variables, eg. house price vs square footage.
  - Classification: Assign a class to an input, eg sentiment analysis or determining digits of zip code on a letter.

# Regression

- We want to find a line that fits our data well (without over fitting) to give us predictive power.
- I'm assuming some basic familiarity with the technique outside of its use in ML.
  - [http://en.wikipedia.org/wiki/Linear\\_regression](http://en.wikipedia.org/wiki/Linear_regression)
  - Least Squares Estimation in particular.

# (Non)linear Regression

- Least Squares Estimation is often used
  - May remember  $w = (X^T X)^{-1} X^T y$
  - If  $X$  is a large matrix, ie a database containing millions of rows, then the matrix inversion can be intractable.
  - Solution: Use Machine Learning

# (Non)linear Regression

- We start with some training data vector
  - $D = \{D_1, D_2, \dots, D_n\}$  where each example  $D_n$  is an input/output pair  $\langle \mathbf{x}_i, y_i \rangle$ .
  - $\mathbf{x}_i = (x_{i,1}, x_{i,2}, \dots, x_{i,d})$  is an input vector with  $d$  features.
    - Features for a house may be square footage, number of bedrooms, etc.
  - $y_i$  is the known output paired with some  $\mathbf{x}_i$ .
    - House price



# (Non)linear Regression

- We want to optimize a hypothesis function  $h$  with output  $z$ .
  - $h(x) = w_0 + w_1x_1 + w_2x_2 \dots + w_dx_d$
  - $w_i$  is a weight
- As the algorithm is shown examples it will produce a predicted output  $z$  which is compared to the expected output  $y$ . Our goal is to find weights that minimize the error  $(y - h(x))^2$ .
  - Should look familiar (Least Squares)

# Gradient Descent


- We achieve this optimization via gradient descent.
  - We have an error function  $E_n = \sum_{0, \dots, n} (h(\mathbf{x}_n) - y_n)^2$
  - We keep changing out weights in  $h(\mathbf{x})$  until we find the the minimum of  $E$ .
  - We change the weights by calculating the gradient and taking a step in proportion to the negative of that gradient.

# Gradient Descent

- Each iteration the weights are changed according to the gradient and we move down the hill toward a minimum.

Picture:

[http://bryannotes.blogspot.com/2014/11/algorithm-stochastic-gradient\\_4.html](http://bryannotes.blogspot.com/2014/11/algorithm-stochastic-gradient_4.html)



# Gradient Descent Algorithm

- The weight update rule is simple:
  - For each example we update each weight according to:
    - $w_i := w_i - \alpha \cdot d/dw_i(E_i)$  (Note these are partial derivatives)
    - For  $h(x) = w_0 + w_1x$  the updates look like:
      - $W_0 := W_0 - \alpha \cdot E$
      - $W_1 := W_1 - \alpha \cdot E \cdot x$
    - $\alpha$  is the learning rate. It determines the step size. Larger  $\alpha$  results in faster descent, but if it's too large the algorithm fails to converge because it oversteps the minimum and ends up bouncing back and forth over it.
  - Repeat until convergence

# Adding the nonlinearity

- A straight line can't fit every data set well.
  - Adding nonlinearity is easy:
    - $h(x) = w_0 + w_1x^2$
  - This is just a simple example. Can use other nonlinearities as desired.
  - As with any regression over fitting can be a problem, so be careful with how you much nonlinearity you add.

# Other Topics

- Logistic Regression uses gradient descent to determine a function for classification. It is very similar to linear regression except it uses discrete outputs and targets.
- Neural Networks are statistical learning algorithms that are significantly more powerful and advanced architectures used for classification and regression.

# Resources

- Many great resources:
  - Wikipedia
  - Coursera Machine Learning Course:
    - <https://class.coursera.org/ml-006>
  - Andrew Ng's Stanford Lectures:
    - <https://www.youtube.com/course?list=ECA89DCFA6ADACE599>
  - CalTech's ML MOOC:
    - <http://work.caltech.edu/previous.html>